

# Feed-Forward Neural Network using Keras and TensorFlow

Muhammad Tamjid Rahman

## Contents

|  |    |
|--|----|
| Loading R packages                                       | 1  |
| 1 Feed-Forward Neural Network using Keras and TensorFlow | 1  |
| 2 A simple neural network                                | 23 |

## Loading R packages

```
library(uuml)
library(keras)
library(tensorflow)
library(tidyverse)
```

## 1 Feed-Forward Neural Network using Keras and TensorFlow

```
mnist <- dataset_mnist()

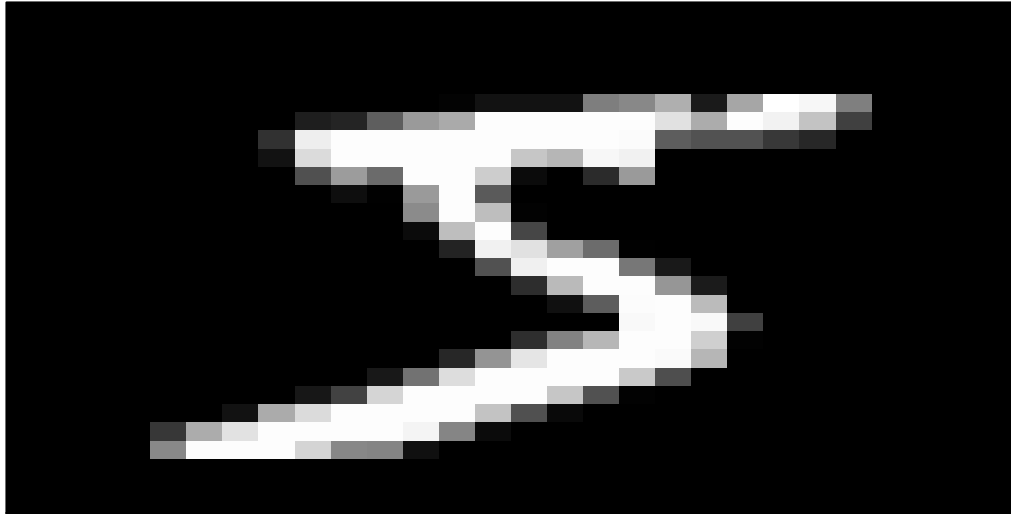
## Loaded Tensorflow version 2.7.0

mnist$train$x <- mnist$train$x/255
mnist$test$x <- mnist$test$x/255
```

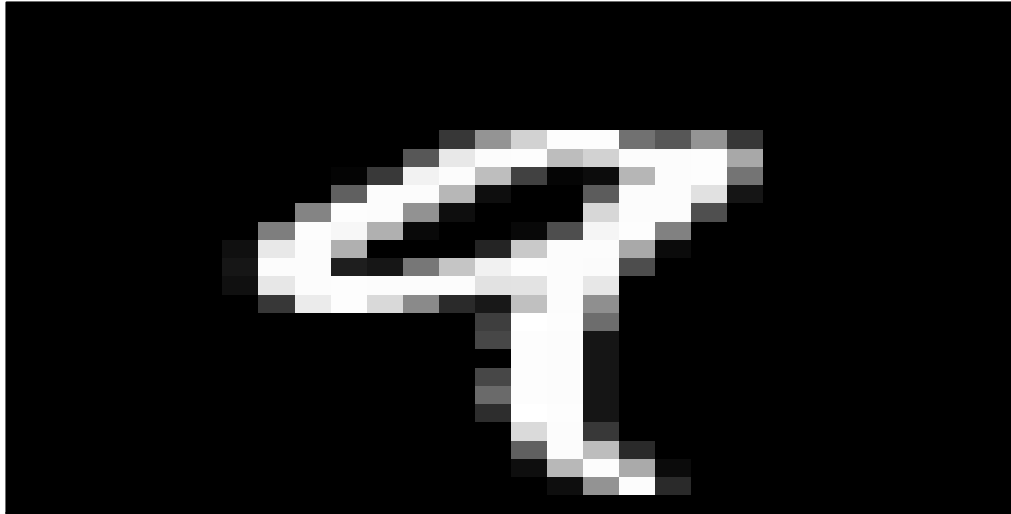
### 1 Visualizing digits

```
idx <- 1
im <- mnist$train$x[idx,,]
# Transpose the image
im <- t(apply(im, 2, rev))
image(1:28, 1:28, im, col=gray((0:255)/255), xlab = "", ylab = "",
xaxt='n', yaxt='n', main=paste(mnist$train$y[idx]))
```

5



```
idx <- 5
im <- mnist$train$x[idx,,]
# Transpose the image
im <- t(apply(im, 2, rev))
image(1:28, 1:28, im, col=gray((0:255)/255), xlab = "", ylab = "",
xaxt='n', yaxt='n', main=paste(mnist$train$y[idx]))
```



## 2 Training and test set size

### Training set

```
dim(as.data.frame( mnist$train))
```

```
## [1] 60000  785
```

```
object.size(mnist$train)
```

```
## 376560792 bytes
```

The training set has 60000 observations with 785 variables and it is 376560792 bytes.

### Test set

```
dim(as.data.frame( mnist$test))
```

```
## [1] 10000  785
```

```
object.size(mnist$test)
```

```
## 62760792 bytes
```

The test set has 10000 observations with 785 variables and it is 62760792 bytes.

### 3 Implementing a feed-forward neural network

a

```
model <- keras_model_sequential() %>%
  layer_flatten(input_shape = c(28, 28)) %>%
  layer_dense(units = 16, activation = "sigmoid") %>%
  layer_dense(10, activation = "softmax")

model %>%
  compile(
    loss = "sparse_categorical_crossentropy",
    optimizer = "adam",
    metrics = "accuracy"
  )

model %>%
  fit(
    x = mnist$train$x, y = mnist$train$y,
    epochs = 5,
    validation_split = 0.3,
    verbose = 2
  )
```

b

```
summary(model)
```

```
## Model: "sequential"
## -----
## Layer (type)                Output Shape          Param #
## =====
## flatten (Flatten)           (None, 784)           0
##
## dense_1 (Dense)              (None, 16)            12560
##
## dense (Dense)                (None, 10)            170
##
## =====
## Total params: 12,730
## Trainable params: 12,730
## Non-trainable params: 0
## -----
```

There are total 12,730 parameters.

c

Layer(type) flatten\_5 (Flatten) is the input layer and it has no parameter.

d

Layer(type) dense\_10 (Dense) is the output layer and it has 170 parameters.

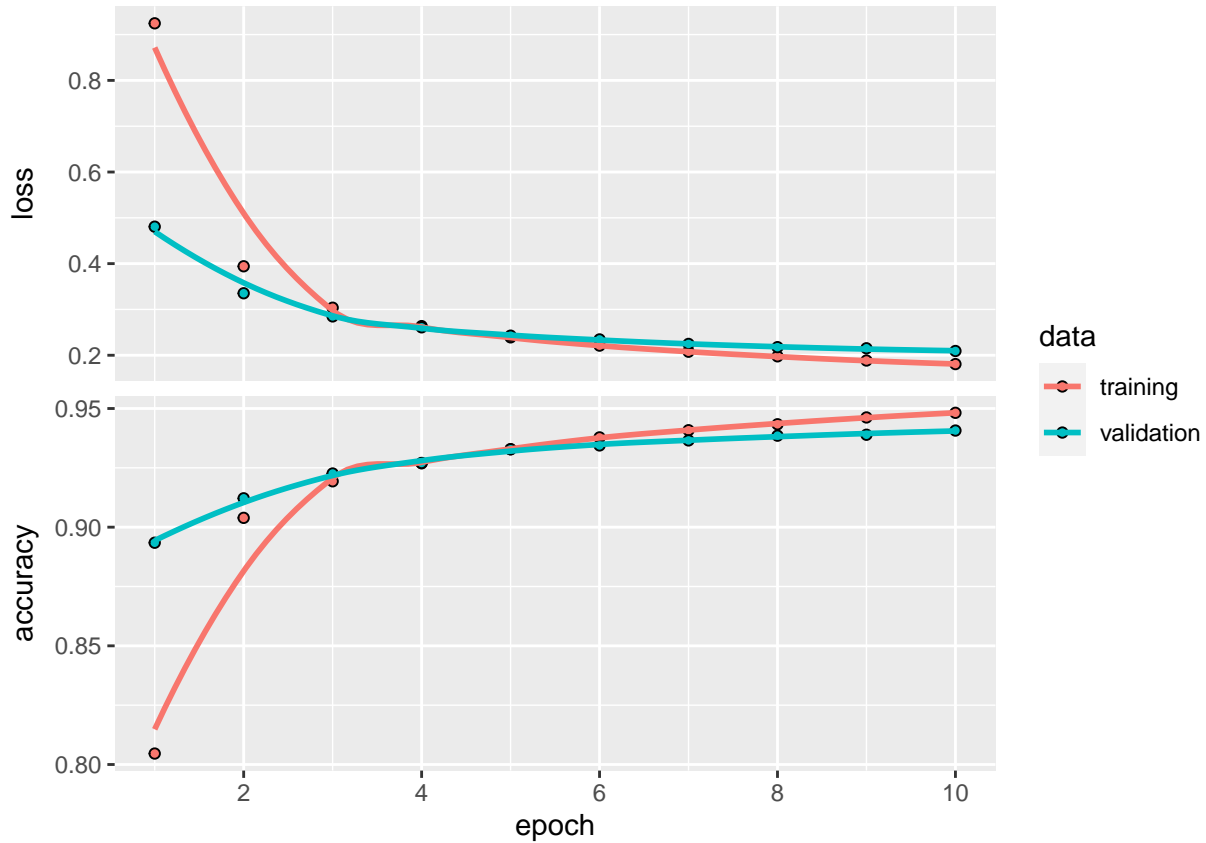
e

```
model %>%
  evaluate(mnist$test$x, mnist$test$y)#, verbose = 0)
```

```
##      loss  accuracy
## 0.2455706 0.9287000
```

The classification accuracy is 93%.

```
## `geom_smooth()` using formula 'y ~ x'
```



## Confusion Matrix

```
pred<- model %>%
  predict((mnist$test$x))%>% k_argmax()
cm<-table(Predicted=as.numeric(pred), Actual=as.numeric(mnist$test$y))
cm
```

```
##      Actual
## Predicted  0   1   2   3   4   5   6   7   8   9
##      0  959   0  12   2   1  12  11   3   5  10
##      1   0 1108   1   2   1   4   3   9   5   6
##      2   0   4  950  14   8   2   8  22   5   1
##      3   2   0   6  923   0  22   1   7  10  10
##      4   1   0  11   1  928   9   8   8   5  26
##      5   8   1   1  17   0  802   9   0  12   4
```

```
##      6      6      3      7      1      5      11  912      0      6      0
##      7      2      2     10     16      4      4      1  958      6      9
##      8      2     17     31     26      5     21      5      1  912      9
##      9      0      0      3      8     30      5      0     20      8  934
```

```
FN=function(mat){
  mat_low_i=list()
  for (i in 1:(nrow(mat)-1)) {

    mat_low_i[i]=list(mat[row(mat) == (col(mat) + i)])
  }

  mat_low_i

  FN1=c()
  for (i in 1:(nrow(mat)-1)) {

    FN1[i]=sum(mat_low_i[[i]])
  }

  p=sum(FN1)

  return(FN=p)
}
```

```
FP=function(mat){
  mat_low_i=list()
  for (i in 1:(nrow(mat)-1)) {

    mat_low_i[i]=list(mat[row(mat) == (col(mat) - i)])
  }

  mat_low_i

  FP1=c()
  for (j in 1:(nrow(mat)-1)) {

    FP1[j]=sum(mat_low_i[[j]])
  }

  p=sum(FP1)

  return(FP=p)
}
```

```
FP=FP(cm)
FP
```

```
## [1] 308
```

```
FN=FN(cm)
FN
```

```
## [1] 306
```

```
TPTN=sum(diag(cm))
TPTN
```

```
## [1] 9386
```

```
accuracy=(TPTN)/(TPTN+FP+FN)
accuracy
```

```
## [1] 0.9386
```

```
precision=TPTN/(TPTN+FP)
precision
```

```
## [1] 0.9682278
```

```
recall=TPTN/(TPTN+FN)
recall
```

```
## [1] 0.9684276
```

| Accuracy | Precision | Recall |
|----------|-----------|--------|
| 93%      | 96%       | 97%    |

4

a

```
model <- keras_model_sequential() %>%
  layer_flatten(input_shape = c(28, 28)) %>%
  layer_dense(units = 128, activation = "sigmoid") %>%
  layer_dense(10, activation = "softmax")
summary(model)
```

```
## Model: "sequential_2"
```

```
## -----
## Layer (type)                Output Shape          Param #
## =====
## flatten_2 (Flatten)         (None, 784)           0
##
## dense_5 (Dense)              (None, 128)           100480
##
## dense_4 (Dense)              (None, 10)            1290
##
## =====
## Total params: 101,770
## Trainable params: 101,770
## Non-trainable params: 0
## -----
```

```
model %>%
  compile(
    loss = "sparse_categorical_crossentropy",
    optimizer = "adam",
    metrics = "accuracy"
  )
```

```

history<-model %>%
  fit(
    x = mnist$train$x, y = mnist$train$y,
    epochs = 10,
    validation_split = 0.3,
    verbose = 2
  )

model %>%
  evaluate(mnist$test$x, mnist$test$y, verbose = 0)

```

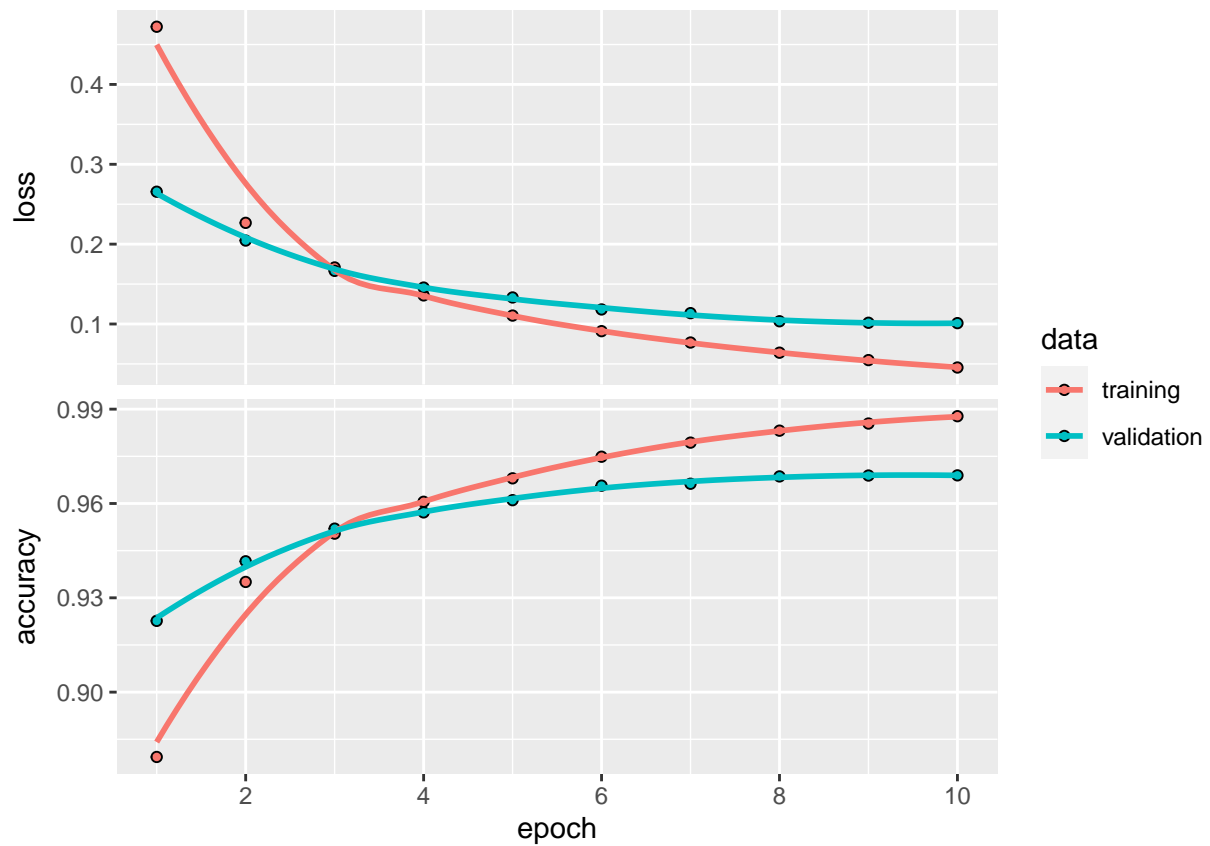
```

##      loss accuracy
## 0.0919233 0.9725000

```

```
plot(history)
```

```
## `geom_smooth()` using formula 'y ~ x'
```



b

```

model <- keras_model_sequential() %>%
  layer_flatten(input_shape = c(28, 28)) %>%
  layer_dense(units = 128, activation = "relu") %>%
  layer_dense(10, activation = "softmax")
summary(model)

```



```
## Model: "sequential_3"
## -----
## Layer (type)                Output Shape          Param #
## =====
## flatten_3 (Flatten)         (None, 784)           0
##
## dense_7 (Dense)              (None, 128)          100480
##
## dense_6 (Dense)              (None, 10)           1290
##
## =====
## Total params: 101,770
## Trainable params: 101,770
## Non-trainable params: 0
## -----
```

```
model %>%
  compile(
    loss = "sparse_categorical_crossentropy",
    optimizer = "adam",
    metrics = "accuracy"
  )

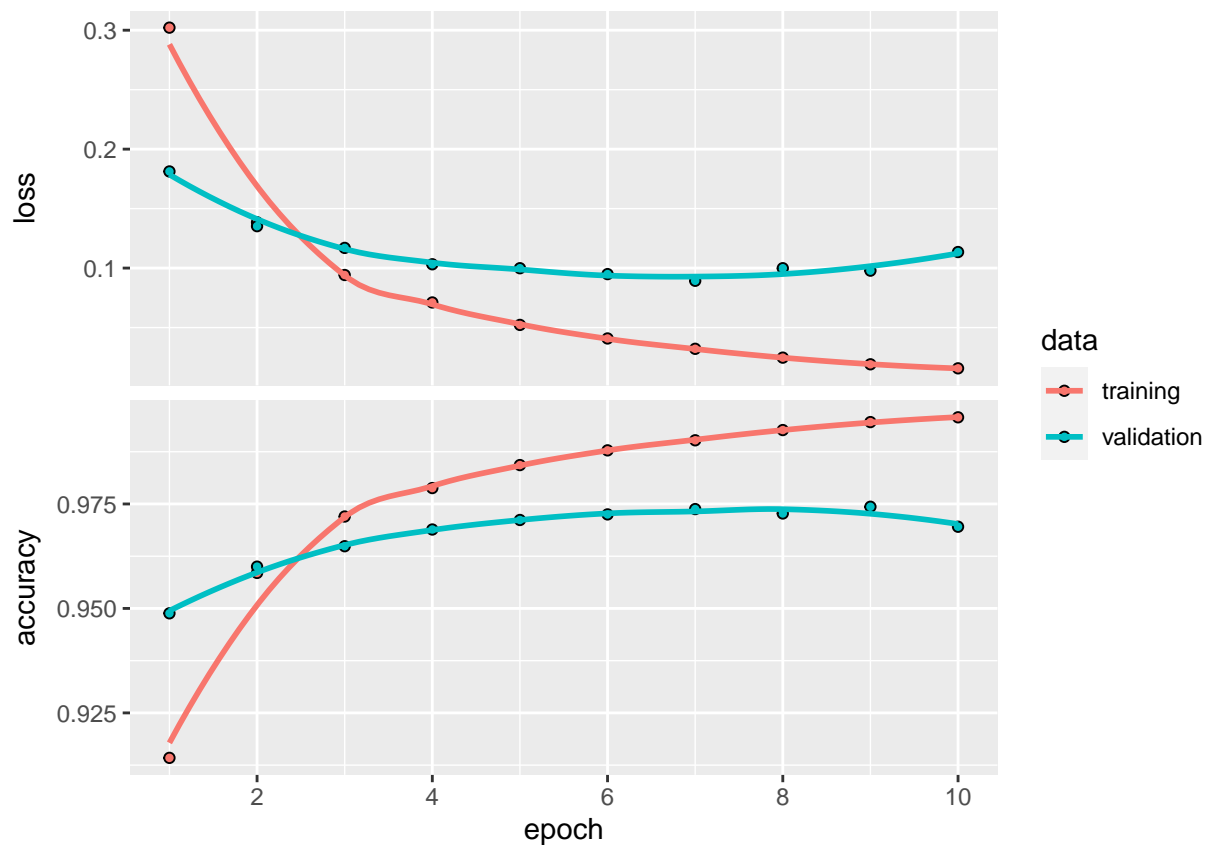
history<-model %>%
  fit(
    x = mnist$train$x, y = mnist$train$y,
    epochs = 10,
    validation_split = 0.3,
    verbose = 2
  )

model %>%
  evaluate(mnist$test$x, mnist$test$y, verbose = 0)
```

```
##      loss  accuracy
## 0.0899054 0.9753000
```

```
plot(history)
```

```
## `geom_smooth()` using formula 'y ~ x'
```



c

```
model <- keras_model_sequential() %>%
  layer_flatten(input_shape = c(28, 28)) %>%
  layer_dense(units = 128, activation = "relu") %>%
  layer_dense(10, activation = "softmax")
summary(model)
```

```
## Model: "sequential_4"
```

```
## -----
## Layer (type)                Output Shape          Param #
## -----
## flatten_4 (Flatten)         (None, 784)           0
##
## dense_9 (Dense)              (None, 128)           100480
##
## dense_8 (Dense)              (None, 10)            1290
##
## =====
## Total params: 101,770
## Trainable params: 101,770
## Non-trainable params: 0
## -----
```

```
model %>%
  compile(
```

```

    loss = "sparse_categorical_crossentropy",
    optimizer = "RMSprop",
    metrics = "accuracy"
)

history<-model %>%
  fit(
    x = mnist$train$x, y = mnist$train$y,
    epochs = 10,
    validation_split = 0.3,
    verbose = 2
  )

model %>%
  evaluate(mnist$test$x, mnist$test$y, verbose = 0)

```

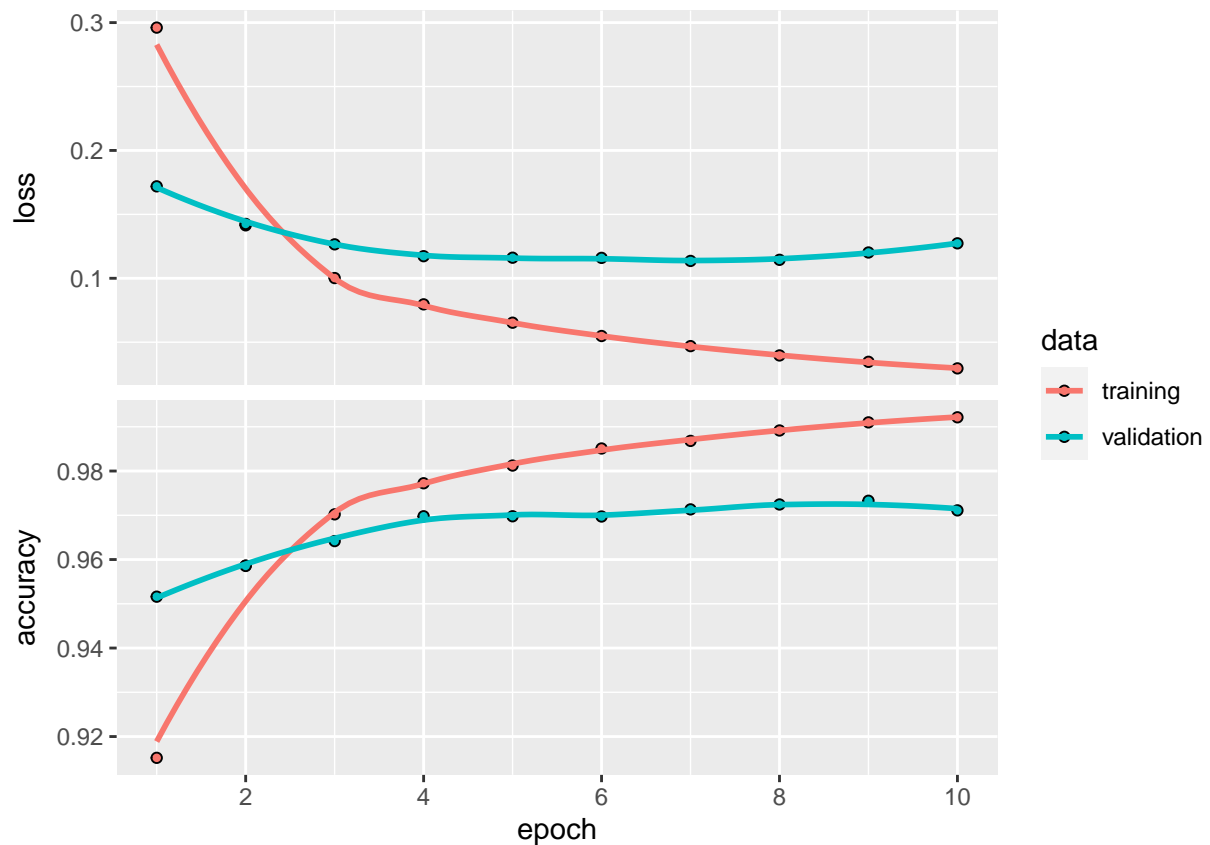
```

##      loss accuracy
## 0.115161 0.975300

```

```
plot(history)
```

```
## `geom_smooth()` using formula 'y ~ x'
```



```

## d
model <- keras_model_sequential() %>%
  layer_flatten(input_shape = c(28, 28)) %>%
  layer_dense(units = 128, activation = "relu") %>%

```

```
layer_dense(10, activation = "softmax")
summary(model)
```

```
## Model: "sequential_5"
## -----
## Layer (type)                Output Shape          Param #
## =====
## flatten_5 (Flatten)         (None, 784)           0
##
## dense_11 (Dense)            (None, 128)           100480
##
## dense_10 (Dense)            (None, 10)            1290
##
## =====
## Total params: 101,770
## Trainable params: 101,770
## Non-trainable params: 0
## -----
```

```
model %>%
  compile(
    loss = "sparse_categorical_crossentropy",
    optimizer = "RMSprop",
    metrics = "accuracy"
  )

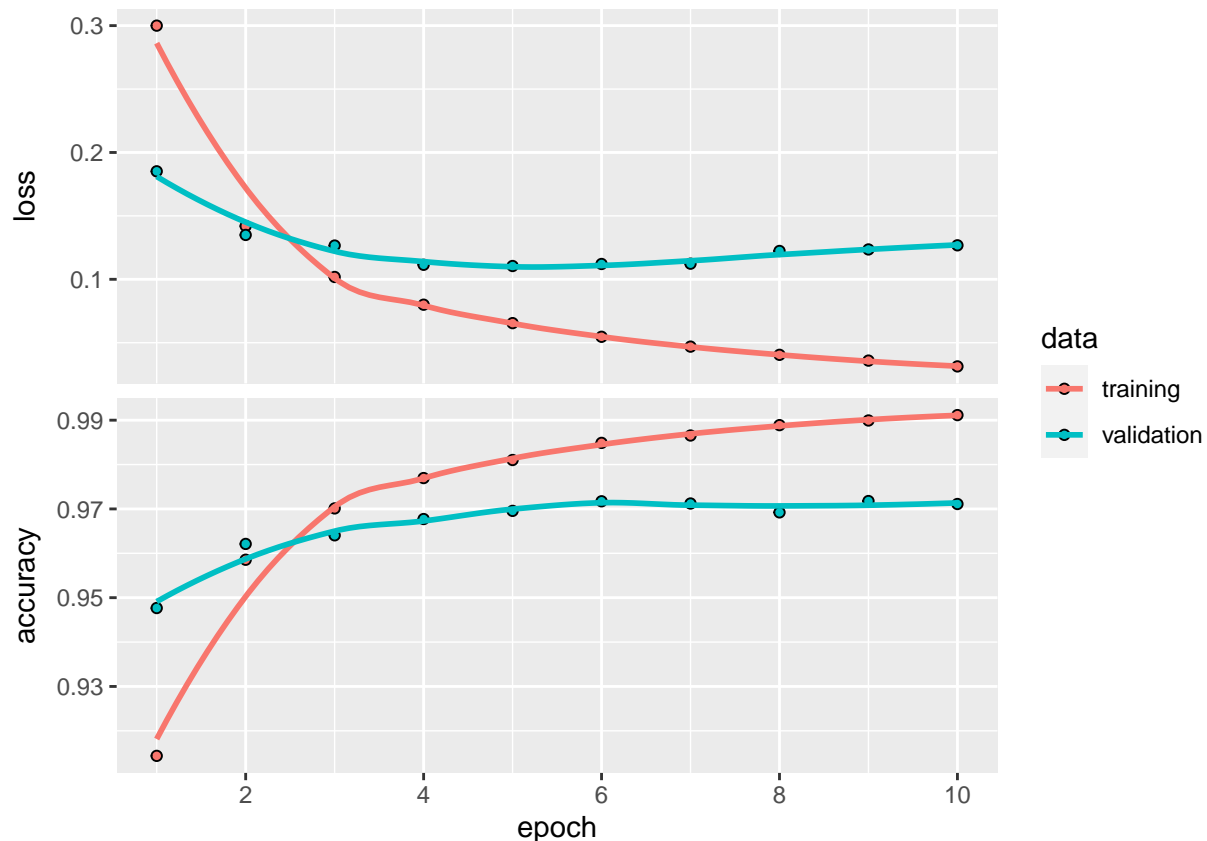
history<-model %>%
  fit(
    x = mnist$train$x, y = mnist$train$y,
    epochs = 10,
    validation_split = 0.3,
    verbose = 2
  )

model %>%
  evaluate(mnist$test$x, mnist$test$y, verbose = 0)
```

```
##      loss  accuracy
## 0.1051387 0.9755000
```

```
plot(history)
```

```
## `geom_smooth()` using formula 'y ~ x'
```



For early stopping regularization I would use kerasR built in function 'EarlyStopping' to avoid being overfitted. This stop training if the improvement(accuracy, loss) of the model for validation has stopped.

e

```
model <- keras_model_sequential() %>%
  layer_flatten(input_shape = c(28, 28)) %>%
  layer_dense(units = 128, activation = "relu") %>%
  layer_dense(units = 128, activation = "relu") %>% #2nd layer
  layer_dense(10, activation = "softmax")
summary(model)
```

```
## Model: "sequential_6"
```

```
## -----
## Layer (type)                Output Shape          Param #
## =====
## flatten_6 (Flatten)         (None, 784)           0
##
## dense_14 (Dense)            (None, 128)           100480
##
## dense_13 (Dense)            (None, 128)           16512
##
## dense_12 (Dense)            (None, 10)            1290
##
## =====
## Total params: 118,282
## Trainable params: 118,282
```

```
## Non-trainable params: 0
## -----
model %>%
  compile(
    loss = "sparse_categorical_crossentropy",
    optimizer = "RMSprop",
    metrics = "accuracy"
  )

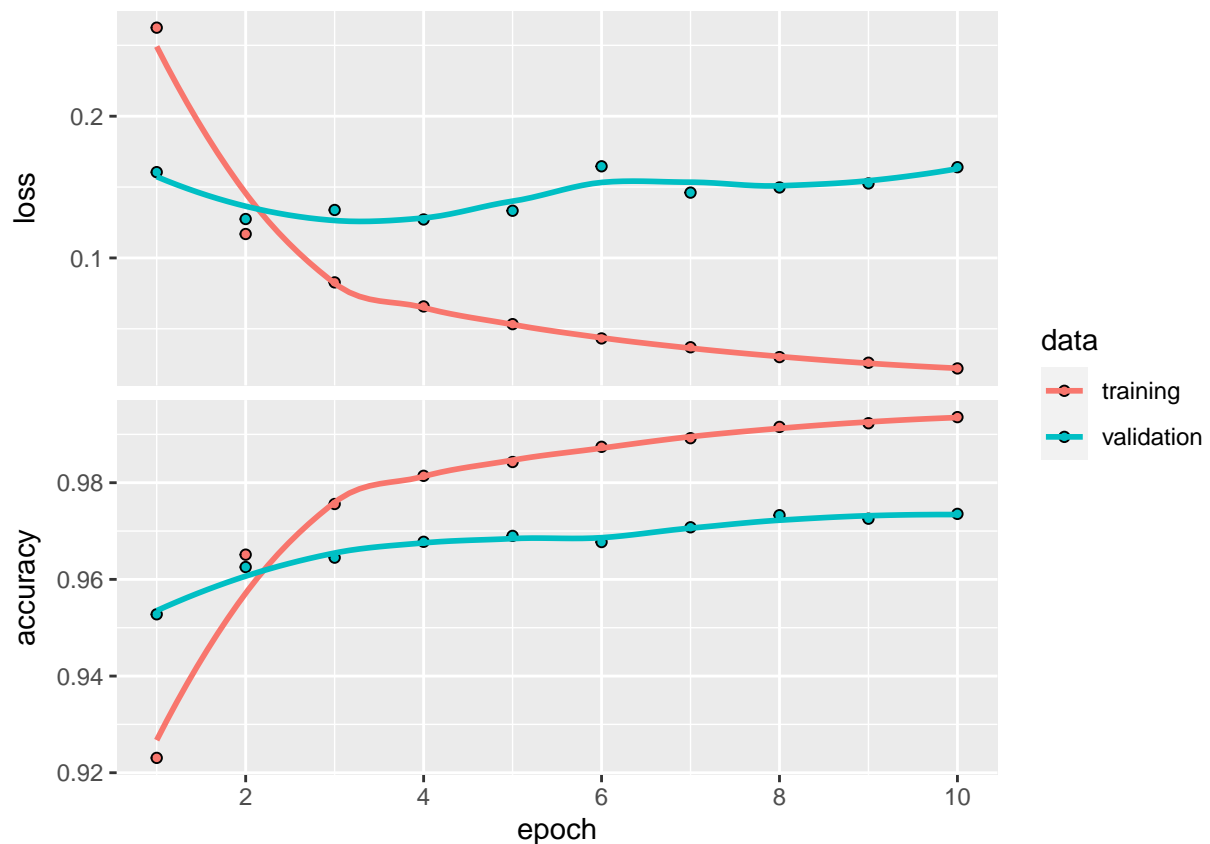
history<-model %>%
  fit(
    x = mnist$train$x, y = mnist$train$y,
    epochs = 10,
    validation_split = 0.3,
    verbose = 2
  )

model %>%
  evaluate(mnist$test$x, mnist$test$y, verbose = 0)

##      loss accuracy
## 0.1210637 0.9785000

plot(history)

## `geom_smooth()` using formula 'y ~ x'
```



There are total 118,282 parameters in the model.

f

```
model <- keras_model_sequential() %>%
  layer_flatten(input_shape = c(28, 28)) %>%
  layer_dense(units = 128, activation = "relu") %>%
  layer_dropout(0.2) %>%
  layer_dense(units = 128, activation = "relu") %>%
  layer_dropout(0.5) %>%
  layer_dense(10, activation = "softmax")
summary(model)
```

```
## Model: "sequential_7"
## -----
## Layer (type)                Output Shape          Param #
## =====
## flatten_7 (Flatten)         (None, 784)           0
##
## dense_17 (Dense)            (None, 128)           100480
##
## dropout_1 (Dropout)         (None, 128)           0
##
## dense_16 (Dense)            (None, 128)           16512
##
## dropout (Dropout)           (None, 128)           0
##
## dense_15 (Dense)            (None, 10)            1290
##
## =====
## Total params: 118,282
## Trainable params: 118,282
## Non-trainable params: 0
## -----
```

```
model %>%
  compile(
    loss = "sparse_categorical_crossentropy",
    optimizer = "RMSprop",
    metrics = "accuracy"
  )

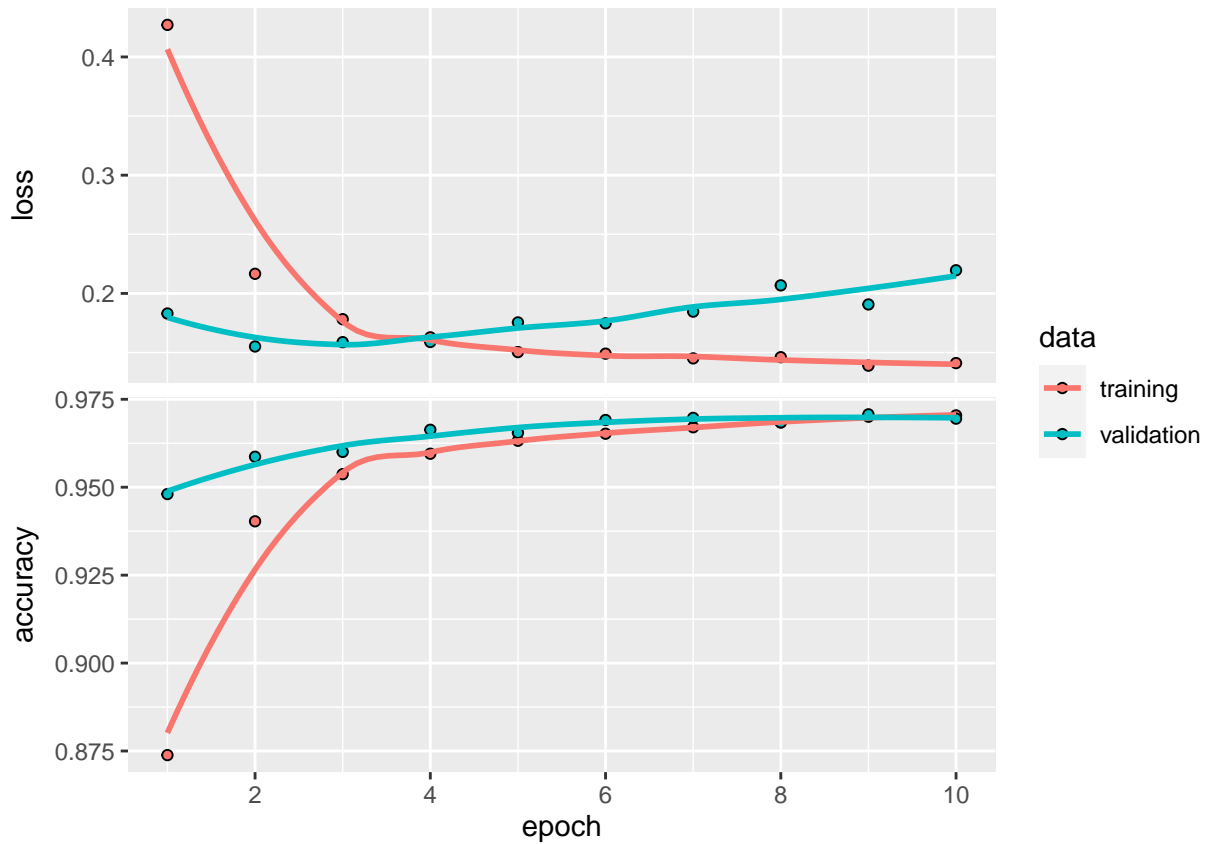
history<-model %>%
  fit(
    x = mnist$train$x, y = mnist$train$y,
    epochs = 10,
    validation_split = 0.3,
    verbose = 2
  )

model %>%
  evaluate(mnist$test$x, mnist$test$y, verbose = 0)

##      loss    accuracy
## 0.1797494 0.9724000
```

```
plot(history)
```

```
## `geom_smooth()` using formula 'y ~ x'
```



```
g
```

```
model <- keras_model_sequential() %>%
  layer_flatten(input_shape = c(28, 28)) %>%
  layer_dense(units = 128, activation = "relu" ) %>%
  layer_batch_normalization() %>%

  layer_dense(units = 128, activation = "relu") %>%
  layer_batch_normalization() %>%

  layer_dense(10, activation = "softmax")
summary(model)
```

```
## Model: "sequential_8"
```

```
## -----
## Layer (type)                Output Shape          Param #
## =====
## flatten_8 (Flatten)         (None, 784)           0
##
## dense_20 (Dense)            (None, 128)           100480
##
## batch_normalization_1 (BatchNormal (None, 128)           512
```



```

## ization)
##
## dense_19 (Dense)                (None, 128)                16512
##
## batch_normalization (BatchNormaliz (None, 128)                512
## ation)
##
## dense_18 (Dense)                (None, 10)                1290
##
## =====
## Total params: 119,306
## Trainable params: 118,794
## Non-trainable params: 512
## -----
model %>%
  compile(
    loss = "sparse_categorical_crossentropy",
    optimizer = "RMSprop",
    metrics = "accuracy"
  )

history<-model %>%
  fit(
    x = mnist$train$x, y = mnist$train$y,
    epochs = 10,
    validation_split = 0.3,
    verbose = 2
  )

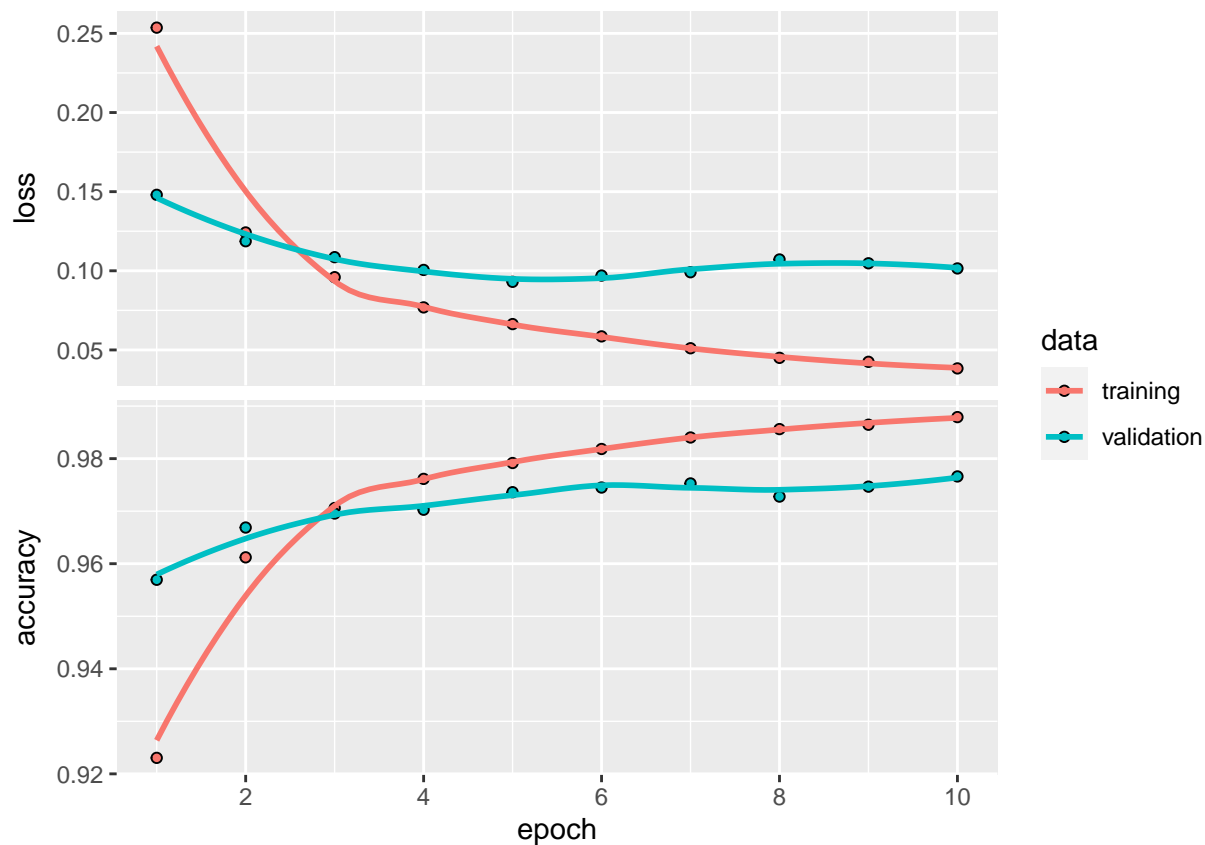
model %>%
  evaluate(mnist$test$x, mnist$test$y, verbose = 0)

##      loss      accuracy
## 0.09347585 0.97570002

plot(history)

## `geom_smooth()` using formula 'y ~ x'

```



## 5 My model

```
model <- keras_model_sequential() %>%
  layer_flatten(input_shape = c(28, 28)) %>%
  layer_dropout(0.2) %>%
  layer_dense(units = 128, activation = "relu" ) %>%
  layer_dropout(0.5) %>%
  layer_batch_normalization() %>%
  layer_dense(units = 256, activation = "relu") %>%
  layer_batch_normalization() %>%
  layer_dense(units = 128, activation = "relu") %>%

  layer_dense(10, activation = "softmax")
summary(model)
```

```
## Model: "sequential_9"
```

```
## -----
## Layer (type)                Output Shape          Param #
## =====
## flatten_9 (Flatten)         (None, 784)           0
##
## dropout_3 (Dropout)         (None, 784)           0
##
## dense_24 (Dense)            (None, 128)          100480
##
```

```
## dropout_2 (Dropout) (None, 128) 0
##
## batch_normalization_3 (BatchNormal (None, 128) 512
## ization)
##
## dense_23 (Dense) (None, 256) 33024
##
## batch_normalization_2 (BatchNormal (None, 256) 1024
## ization)
##
## dense_22 (Dense) (None, 128) 32896
##
## dense_21 (Dense) (None, 10) 1290
##
## =====
## Total params: 169,226
## Trainable params: 168,458
## Non-trainable params: 768
## -----
```

```
model %>%
  compile(
    loss = "sparse_categorical_crossentropy",
    optimizer = "adam",
    metrics = "accuracy"
  )

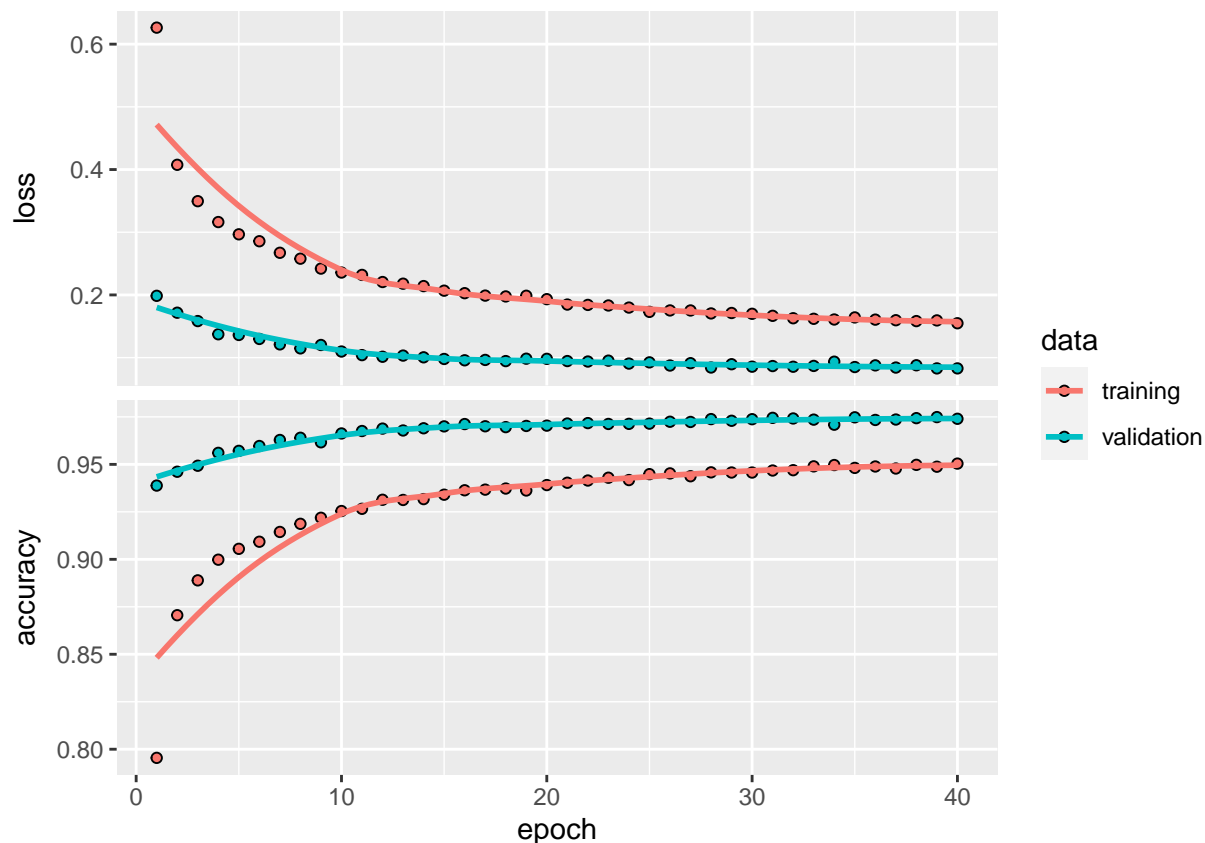
history<-model %>%
  fit(
    x = mnist$train$x, y = mnist$train$y,
    epochs = 40,
    validation_split = 0.3,
    verbose = 2
  )

model %>%
  evaluate(mnist$test$x, mnist$test$y, verbose = 0)
```

```
##      loss      accuracy
## 0.07251313 0.97970003
```

```
plot(history)
```

```
## `geom_smooth()` using formula 'y ~ x'
```



In my model, 3 hidden layer were used with the same activation function reLU. Batch normalized and dropout of layer were used in 2 hidden layers. After 30 epochs accuracy and loss seem to be stabled. The accuracy is 98%. Validation set accuracy is close to 95%. So, the model is well fitted, not overfitted.

## 6 Two digits that the network has classified incorrectly

```
library(tidyverse)

pred<- model %>%
  predict((mnist$test$x))%>% k_argmax()

df<-cbind(1:length(mnist$test$y),mnist$test$y,as.numeric(pred))

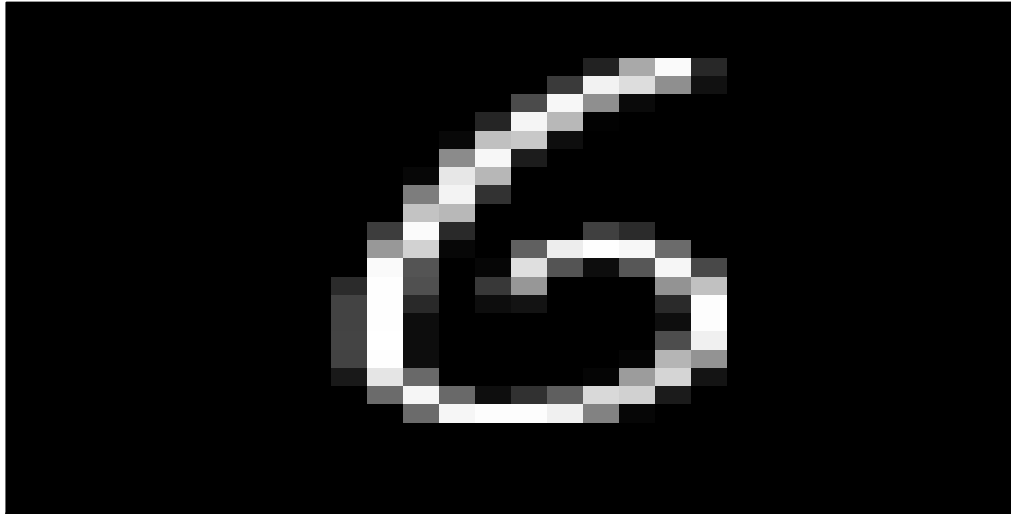
subset(df, df[,2] !=df[,3])[1:2,]

##      [,1] [,2] [,3]
## [1,]   19    3    8
## [2,]  152    9    8
```

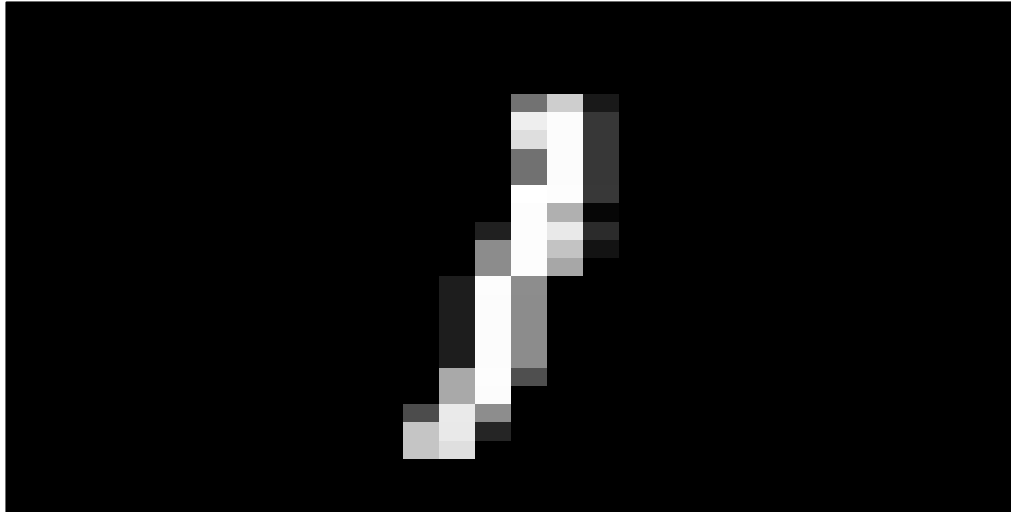
Image with index 19, 125 were miss-classified. 19th image is 3 but predicted as 8 and 19th image is 7 but predicted as 4.

```
idx <- 19
im <- mnist$train$x[idx,,]
# Transpose the image
im <- t(apply(im, 2, rev))
image(1:28, 1:28, im, col=gray((0:255)/255), xlab = "", ylab = "",
xaxt='n', yaxt='n', main=paste(mnist$test$y[idx]))
```

3



```
idx <- 125
im <- mnist$train$x[idx,,]
# Transpose the image
im <- t(apply(im, 2, rev))
image(1:28, 1:28, im, col=gray((0:255)/255), xlab = "", ylab = "",
xaxt='n', yaxt='n', main=paste(mnist$test$y[idx]))
```



After visualizing we see that both 19th and 125th were labeling incorrectly. This can be one of the reasons of miss-classification.

7

## Confusion Matrix

```
cm1<-table(Predicted=as.numeric(pred), Actual=as.numeric(mnist$test$y))
cm1
```

```
##          Actual
## Predicted    0    1    2    3    4    5    6    7    8    9
##          0  973    0    4    0    1    2    5    2    5    1
##          1    0 1132    3    1    0    0    2    8    1    5
##          2    0    0 1009    3    1    0    0    8    4    0
##          3    0    0    2  986    0    6    1    1    5    5
##          4    0    0    3    0  961    1    2    0    4    8
##          5    1    1    0    5    0  872    3    0    3    8
##          6    0    1    1    0    7    6  941    0    1    0
##          7    1    0    7    7    1    1    0 1003    4    4
##          8    2    1    3    5    2    3    4    0  945    3
##          9    3    0    0    3    9    1    0    6    2  975
```

```
FP=FP(cm1)
FP
```

```
## [1] 115
```

```

FN=FN(cm1)
FN

## [1] 88
TPTN=sum(diag(cm1))
TPTN

## [1] 9797
accuracy=(TPTN)/(TPTN+FP+FN)
accuracy

## [1] 0.9797
precision=TPTN/(TPTN+FP)
precision

## [1] 0.9883979
recall=TPTN/(TPTN+FN)
recall

## [1] 0.9910976

```

| Accuracy | Precision | Recall |
|----------|-----------|--------|
| 98%      | 99%       | 99%    |

We get better accuracy than the validation set.

## 2 A simple neural network

A simple neural network function is going to be implemented.

```

mini_net<- function(X, W, c, w, b){
  w1 <- X %*% W + c
  a1 <- matrix(NA,nrow(X),ncol(W ))
  for(i in 1:nrow(X)){
    for(j in 1:ncol(W)){
      a1[i,j] <- max(0,w1[i,j]) # activation function
    }
  }
  return(y1=a1%*%w+b)
}

W <- matrix(1, nrow = 2, ncol = 2)
c <- matrix(c(0, -1), ncol = 2, nrow = 4, byrow = TRUE)
X <- matrix(c(0,0,1,1,0,1,0,1), ncol = 2)
w <- matrix(c(1, -2), ncol = 1)
b <- 0
mini_net(X, W, c, w, b)

##      [,1]
## [1,]    0
## [2,]    1

```

```
## [3,]    1
## [4,]    0

##      [,1]
## [1,]  0.0
## [2,]  0.9
## [3,]  0.9
## [4,]  0.2
```

## 2

Changing the value  $W_{11}$  to 0.

```
W[1,1]=0
mini_net(X, W, c, w, b)
```

```
##      [,1]
## [1,]    0
## [2,]    1
## [3,]    0
## [4,]   -1
```

The above is the result we got.

## 3

For the neural network we used reLU as activation function. So, the output function is  $\text{relu}(X\% * \%W + c)\% * \%w\% + b$ . For simple problem this network might be reasonable but in reality not. It's better to use more complex network. We could have used sigmoid instead of reLU.

## 4 Implementing a mean squared error loss function

```
W[1,1]=1
mini_net_loss=function(y, X, W, c, w, b){
  y1=mini_net(X, W, c, w, b)
  MSE=mean((y1-y)^2)
  return(MSE)
}
y <- c(0,1,1,0)
mini_net_loss(y, X, W, c, w, b)
```

```
## [1] 0
mini_net_loss(y, X, 0.9*W, c, w, b)
```

```
## [1] 0.015
```

## 5

```
W[1,1]=0
mini_net_loss(y, X, W, c, w, b)
```

```
## [1] 0.5
```

The value of the loss function is 0.5 .